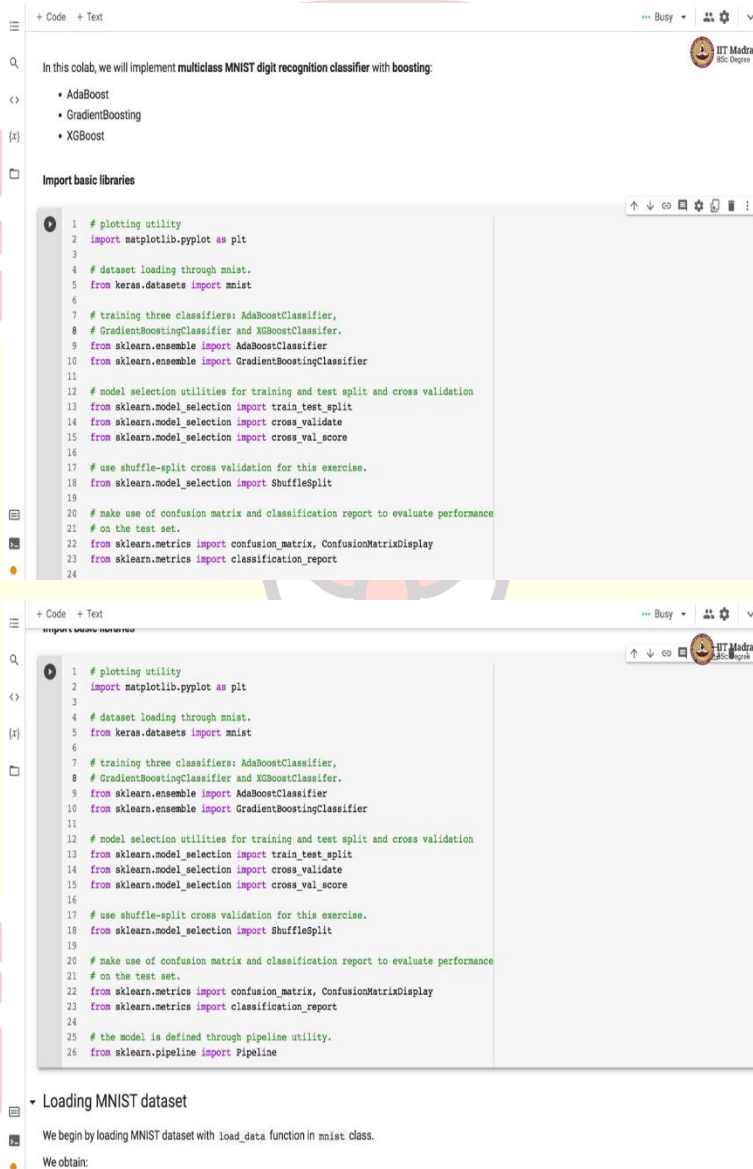


**IIT Madras**  
ONLINE DEGREE

**Machine Learning Practice**  
**Professor. Ashish Tendulkar**  
**Indian Institute of Technology, Madras**  
**AdaBoost and GradientBoost Classifier on MNIST**

**(Refer Slide Time: 00:10)**



The image displays two screenshots of a Jupyter Notebook interface. The top screenshot shows the notebook's title bar with 'Code' and 'Text' tabs, a 'Busy' status indicator, and the IIT Madras logo. Below the title bar, a text cell contains the introductory text: 'In this colab, we will implement multiclass MNIST digit recognition classifier with boosting:'. This is followed by a bulleted list of classifiers: 'AdaBoost', 'GradientBoosting', and 'XGBoost'. Below the list, a code cell is shown with the following Python code:

```
1 # plotting utility
2 import matplotlib.pyplot as plt
3
4 # dataset loading through mnist.
5 from keras.datasets import mnist
6
7 # training three classifiers: AdaBoostClassifier,
8 # GradientBoostingClassifier and XGBoostClassifier.
9 from sklearn.ensemble import AdaBoostClassifier
10 from sklearn.ensemble import GradientBoostingClassifier
11
12 # model selection utilities for training and test split and cross validation
13 from sklearn.model_selection import train_test_split
14 from sklearn.model_selection import cross_validate
15 from sklearn.model_selection import cross_val_score
16
17 # use shuffle-split cross validation for this exercise.
18 from sklearn.model_selection import ShuffleSplit
19
20 # make use of confusion matrix and classification report to evaluate performance
21 # on the test set.
22 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
23 from sklearn.metrics import classification_report
24
```

The bottom screenshot shows the same notebook with the code cell expanded to line 26, where the Pipeline class is imported:

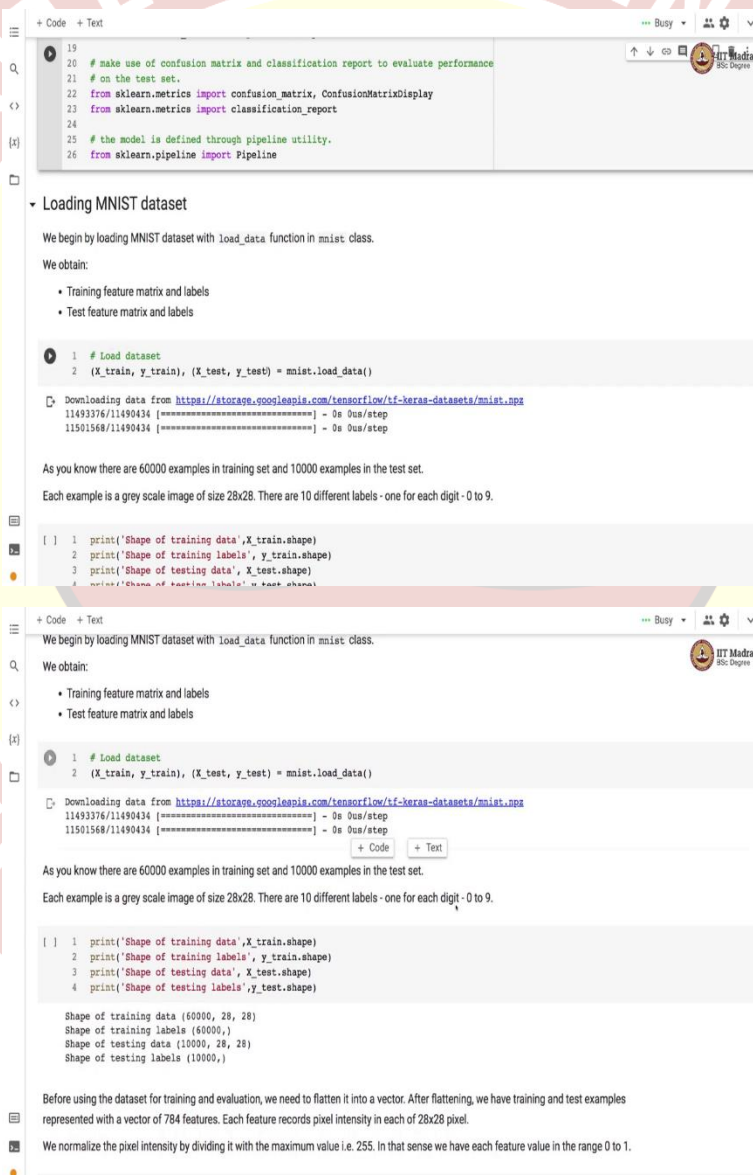
```
26 from sklearn.pipeline import Pipeline
```

Below the code cell, a section titled 'Loading MNIST dataset' is visible, containing the text: 'We begin by loading MNIST dataset with load\_data function in mnist class. We obtain:'.

Namaste! Welcome to the next video of Machine Learning Practice Course. In this video, we will implement multiclass MNIST digit recognition classifier with boosting, we will be using 3 boosting classifiers, AdaBoostClassifier, GradientBoostingClassifier and XGBoostclassifier. We begin by importing our usual Python libraries. We will be making use of matplotlib.pyplot for plotting.

We will be loading dataset through MNIST, we will be training 3 classifiers, AdaBoostClassifier and GradientBoostingClassifier that are implemented as part of sklearn.ensemble module. Then we have a bunch of model selection utilities like train \_test \_split, and cross-validation utilities. We will be using ShuffleSplit cross-validation for this exercise. We will make use of confusion \_matrix and classification \_report to evaluate the performance on the test set. And the model is defined through the pipeline utility.

(Refer Slide Time: 01:14)



```
19 # make use of confusion matrix and classification report to evaluate performance
20 # on the test set.
21 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
22 from sklearn.metrics import classification_report
23
24 # the model is defined through pipeline utility.
25 from sklearn.pipeline import Pipeline
```

▼ Loading MNIST dataset

We begin by loading MNIST dataset with load\_data function in mnist class.

We obtain:

- Training feature matrix and labels
- Test feature matrix and labels

```
1 # Load dataset
2 (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.zip>

11493376/11490434 [=====] - 0s 0us/step

11501568/11490434 [=====] - 0s 0us/step

As you know there are 60000 examples in training set and 10000 examples in the test set.

Each example is a grey scale image of size 28x28. There are 10 different labels - one for each digit - 0 to 9.

```
1 print('Shape of training data', X_train.shape)
2 print('Shape of training labels', y_train.shape)
3 print('Shape of testing data', X_test.shape)
4 print('Shape of testing labels', y_test.shape)
```

We begin by loading MNIST dataset with load\_data function in mnist class.

We obtain:

- Training feature matrix and labels
- Test feature matrix and labels

```
1 # Load dataset
2 (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.zip>

11493376/11490434 [=====] - 0s 0us/step

11501568/11490434 [=====] - 0s 0us/step

As you know there are 60000 examples in training set and 10000 examples in the test set.

Each example is a grey scale image of size 28x28. There are 10 different labels - one for each digit - 0 to 9.

```
1 print('Shape of training data', X_train.shape)
2 print('Shape of training labels', y_train.shape)
3 print('Shape of testing data', X_test.shape)
4 print('Shape of testing labels', y_test.shape)
```

Shape of training data (60000, 28, 28)

Shape of training labels (60000,)

Shape of testing data (10000, 28, 28)

Shape of testing labels (10000,)

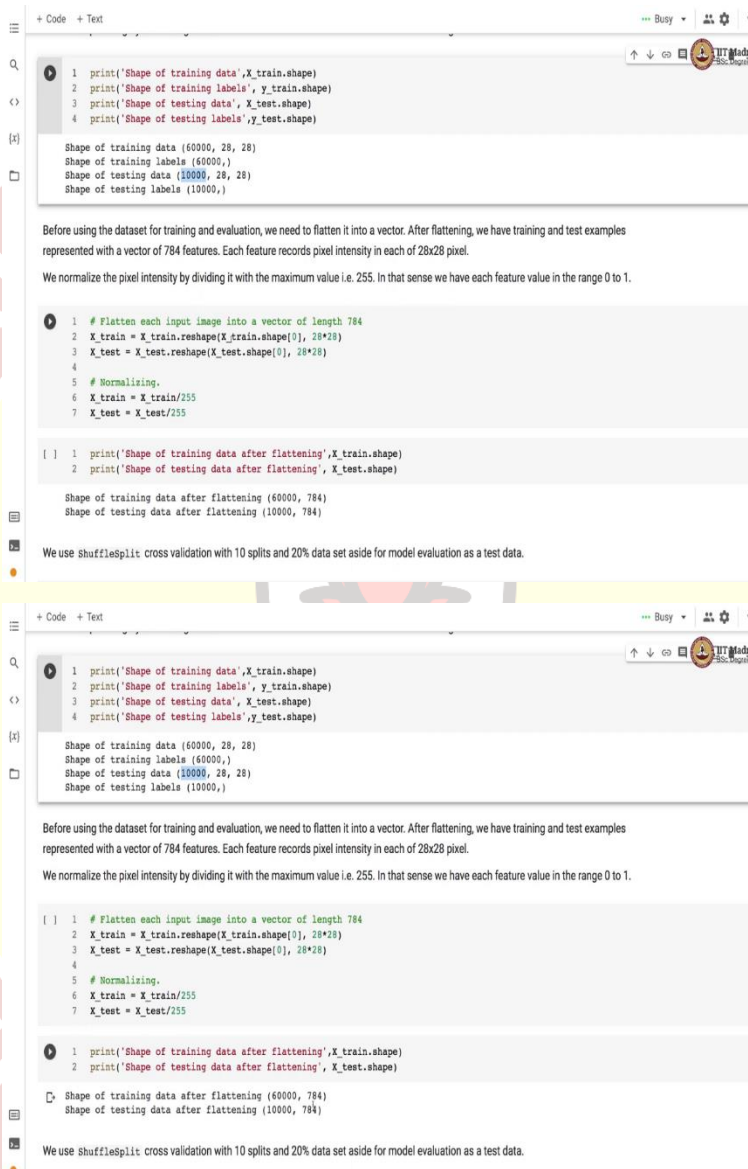
Before using the dataset for training and evaluation, we need to flatten it into a vector. After flattening, we have training and test examples represented with a vector of 784 features. Each feature records pixel intensity in each of 28x28 pixel.

We normalize the pixel intensity by dividing it with the maximum value i.e. 255. In that sense we have each feature value in the range 0 to 1.

We begin by loading MNIST dataset with load \_data function in MNIST class, we obtain training feature matrix and labels as well as test feature matrix and labels. As you know, there are 60,000

examples in the training set, and 10,000 examples in the test set. Each example is a grey scale image of size  $28 \times 28$ . And there are 10 different labels 1 for each digit between 0 to 9.

(Refer Slide Time: 01:44)



The image displays two screenshots of a Jupyter Notebook interface, showing the initial steps of data loading and preprocessing for a digit recognition task. The background features a large, semi-transparent watermark of the Indian Institute of Technology Madras logo.

**Top Screenshot:**

```
1 print('Shape of training data',X_train.shape)
2 print('Shape of training labels', y_train.shape)
3 print('Shape of testing data', X_test.shape)
4 print('Shape of testing labels',y_test.shape)
```

Shape of training data (60000, 28, 28)  
Shape of training labels (60000,)  
Shape of testing data (10000, 28, 28)  
Shape of testing labels (10000,)

Before using the dataset for training and evaluation, we need to flatten it into a vector. After flattening, we have training and test examples represented with a vector of 784 features. Each feature records pixel intensity in each of  $28 \times 28$  pixel.

We normalize the pixel intensity by dividing it with the maximum value i.e. 255. In that sense we have each feature value in the range 0 to 1.

```
1 # Flatten each input image into a vector of length 784
2 X_train = X_train.reshape(X_train.shape[0], 28*28)
3 X_test = X_test.reshape(X_test.shape[0], 28*28)
4
5 # Normalizing.
6 X_train = X_train/255
7 X_test = X_test/255
```

```
1 print('Shape of training data after flattening',X_train.shape)
2 print('Shape of testing data after flattening', X_test.shape)
```

Shape of training data after flattening (60000, 784)  
Shape of testing data after flattening (10000, 784)

We use `shuffleSplit`: cross validation with 10 splits and 20% data set aside for model evaluation as a test data.

**Bottom Screenshot:**

```
1 print('Shape of training data',X_train.shape)
2 print('Shape of training labels', y_train.shape)
3 print('Shape of testing data', X_test.shape)
4 print('Shape of testing labels',y_test.shape)
```

Shape of training data (60000, 28, 28)  
Shape of training labels (60000,)  
Shape of testing data (10000, 28, 28)  
Shape of testing labels (10000,)

Before using the dataset for training and evaluation, we need to flatten it into a vector. After flattening, we have training and test examples represented with a vector of 784 features. Each feature records pixel intensity in each of  $28 \times 28$  pixel.

We normalize the pixel intensity by dividing it with the maximum value i.e. 255. In that sense we have each feature value in the range 0 to 1.

```
1 # Flatten each input image into a vector of length 784
2 X_train = X_train.reshape(X_train.shape[0], 28*28)
3 X_test = X_test.reshape(X_test.shape[0], 28*28)
4
5 # Normalizing.
6 X_train = X_train/255
7 X_test = X_test/255
```

```
1 print('Shape of training data after flattening',X_train.shape)
2 print('Shape of testing data after flattening', X_test.shape)
```

Shape of training data after flattening (60000, 784)  
Shape of testing data after flattening (10000, 784)

We use `shuffleSplit`: cross validation with 10 splits and 20% data set aside for model evaluation as a test data.

```
+ Code + Text
3 print('shape of training data', X_train.shape)
4 print('Shape of testing labels', y_test.shape)

Shape of training data (60000, 28, 28)
Shape of training labels (60000,)
Shape of testing data (10000, 28, 28)
Shape of testing labels (10000,)

Before using the dataset for training and evaluation, we need to flatten it into a vector. After flattening, we have training and test examples represented with a vector of 784 features. Each feature records pixel intensity in each of 28x28 pixel.

We normalize the pixel intensity by dividing it with the maximum value i.e. 255. In that sense we have each feature value in the range 0 to 1.

[ ] 1 # Flatten each input image into a vector of length 784
2 X_train = X_train.reshape(X_train.shape[0], 28*28)
3 X_test = X_test.reshape(X_test.shape[0], 28*28)
4
5 # Normalizing.
6 X_train = X_train/255
7 X_test = X_test/255

[ ] 1 print('Shape of training data after flattening', X_train.shape)
2 print('Shape of testing data after flattening', X_test.shape)

Shape of training data after flattening (60000, 784)
Shape of testing data after flattening (10000, 784)

We use shuffleSplit cross validation with 10 splits and 20% data set aside for model evaluation as a test data.

[ ] 1 cv = ShuffleSplit(n_splits=10, test_size=0.2, random_state=42)

We define two functions:
```

Before using the dataset for training and evaluation, we need to flatten it into a vector. After flattening we have training and test example represented with 784 features. Each feature records pixel intensity in each of the  $28 \times 28$  pixels. We normalize the pixel intensity by dividing it with maximum value that is 255. In that sense, we have each feature value in the range between 0 to 1.

**(Refer Slide Time: 02:16)**

```
+ Code + Text
We define two functions:

1. train_classifiers contains a common code for training classifiers for MNIST multiclass classification problem.
   - It takes estimator, feature matrix, labels, cross validation strategy and name of the classifier as input.
   - It first fits the estimator with feature matrix and labels.
   - It obtains cross validated f1_macro score for training set with 10-fold shuffleSplit cross validation and prints it.

[ ] 1 def train_classifiers(estimator, X_train, y_train, cv, name):
2     estimator.fit(X_train, y_train)
3     cv_train_score = cross_val_score(estimator, X_train, y_train,
4                                     cv=cv, scoring='f1_macro')
5     print(f'On an average, {name} model has f1 score of '
6           f'{cv_train_score.mean():.3f} +/- {cv_train_score.std():.3f} on the training set.')

2. The eval function takes estimator, test feature matrix and labels as input and produce classification report and confusion matrix.
   - It first predicts labels for the test set.
   - Then it uses these predicted reports for calculating various evaluation metrics like precision, recall, f1 score and accuracy for each of the 10 classes.
   - It also obtains a confusion matrix by comparing these predictions with the actual labels and displays it with confusionMatrixDisplay utility.

[ ] 1 def eval(estimator, X_test, y_test):
2     y_pred = estimator.predict(X_test)
3
4     print("# Classification report")
5     print(classification_report(y_test, y_pred))
6
7     print("# Confusion matrix")
```



We define two functions:

1. `train_classifier` contains a common code for training classifiers for MNIST multiclass classification problem.
  - It takes `estimator`, feature matrix, labels, cross validation strategy and name of the classifier as input.
  - It first fits the estimator with feature matrix and labels.
  - It obtains cross validated `f1_macro` score for training set with 10-fold `ShuffleSplit` cross validation and prints it.
2. The `eval` function takes estimator, test feature matrix and labels as input and produce classification report and confusion matrix.
  - It first predicts labels for the test set.
  - Then it uses these predicted reports for calculating various evaluation metrics like precision, recall, f1 score and accuracy for each of the 10 classes.
  - It also obtains a confusion matrix by comparing these predictions with the actual labels and displays it with `ConfusionMatrixDisplay` utility.

```

1 def train_classifier(estimator, X_train, y_train, cv, name):
2     estimator.fit(X_train, y_train)
3     cv_train_score = cross_val_score(estimator, X_train, y_train,
4                                     cv=cv, scoring='f1_macro')
5     print(f'On an average, {name} model has f1 score of {cv_train_score.mean():.3f} +/- {cv_train_score.std():.3f} on the training set.')
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100

```

```

1 def eval(estimator, X_test, y_test):
2     y_pred = estimator.predict(X_test)
3
4     print("# Classification report")
5     print(classification_report(y_test, y_pred))
6
7     print("# Confusion matrix")
8     disp = ConfusionMatrixDisplay(
9         confusion_matrix=confusion_matrix(y_test, y_pred))
10    disp.plot()
11    plt.title('Confusion matrix')
12    plt.show()

```

We use `ShuffleSplit` cross-validation strategy with 10-folds. And we set aside 20 %examples as test data for model evaluation. We define 2 functions. 1 is `train_classifier` that contains common code for training classifier for MNIST multiclass classification problem. It takes estimator feature matrix labels, cross-validation strategy and name of the classifier as input. It first fits the estimator with feature matrix and labels. It obtains cross validated `f1_macro` score for training set with 10-fold `ShuffleSplit` cross-validation and it prints it with these 2 statements.

**(Refer Slide Time: 03:03)**

2. The `eval` function takes estimator, test feature matrix and labels as input and produce classification report and confusion matrix.

- It first predicts labels for the test set.
- Then it uses these predicted reports for calculating various evaluation metrics like precision, recall, f1 score and accuracy for each of the 10 classes.
- It also obtains a confusion matrix by comparing these predictions with the actual labels and displays it with `ConfusionMatrixDisplay` utility.

```

1 def eval(estimator, X_test, y_test):
2     y_pred = estimator.predict(X_test)
3
4     print("# Classification report")
5     print(classification_report(y_test, y_pred))
6
7     print("# Confusion matrix")
8     disp = ConfusionMatrixDisplay(
9         confusion_matrix=confusion_matrix(y_test, y_pred))
10    disp.plot()
11    plt.title('Confusion matrix')
12    plt.show()

```

Let's train two classifiers with default parameters.

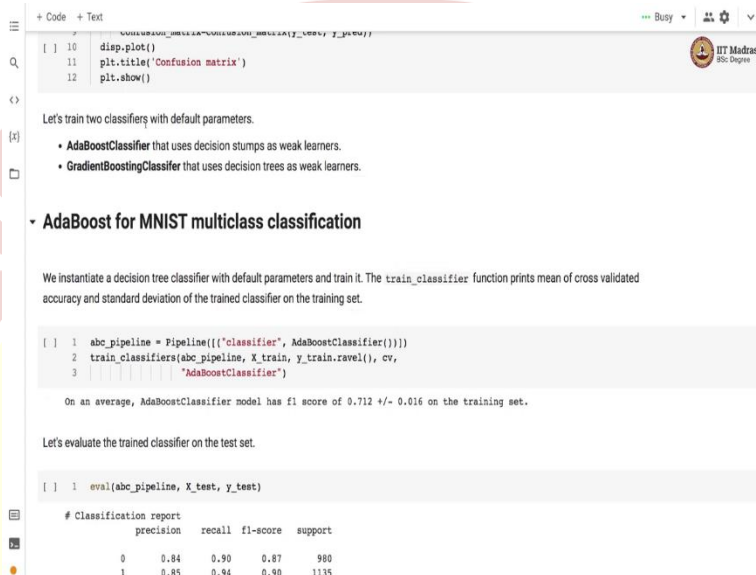
- `AdaBoostClassifier` that uses decision stumps as weak learners.
- `GradientBoostingClassifier` that uses decision trees as weak learners.

AdaBoost for MNIST multiclass classification

The `eval` function on the other hand takes estimator test feature matrix and labels as input and produce classification `_report` and confusion `_matrix`. It first predicts label for the test set. Then it

uses the predicted labels for obtaining the classification \_report. And in classification \_report, we have precision recall, and f1 \_score for each of the 10 classes. It also obtains confusion \_matrix by comparing these predictions and displayed with ConfusionMatrixDisplay utility.

(Refer Slide Time: 03:38)



```
10 disp.plot()
11 plt.title('Confusion matrix')
12 plt.show()
```

Let's train two classifiers with default parameters.

- **AdaBoostClassifier** that uses decision stumps as weak learners.
- **GradientBoostingClassifier** that uses decision trees as weak learners.

• **AdaBoost for MNIST multiclass classification**

We instantiate a decision tree classifier with default parameters and train it. The `train_classifier` function prints mean of cross validated accuracy and standard deviation of the trained classifier on the training set.

```
1 abc_pipeline = Pipeline([("classifier", AdaBoostClassifier())])
2 train_classifiers(abc_pipeline, X_train, y_train.ravel(), cv,
3                  "AdaBoostClassifier")
```

On an average, AdaBoostClassifier model has f1 score of 0.712 +/- 0.016 on the training set.

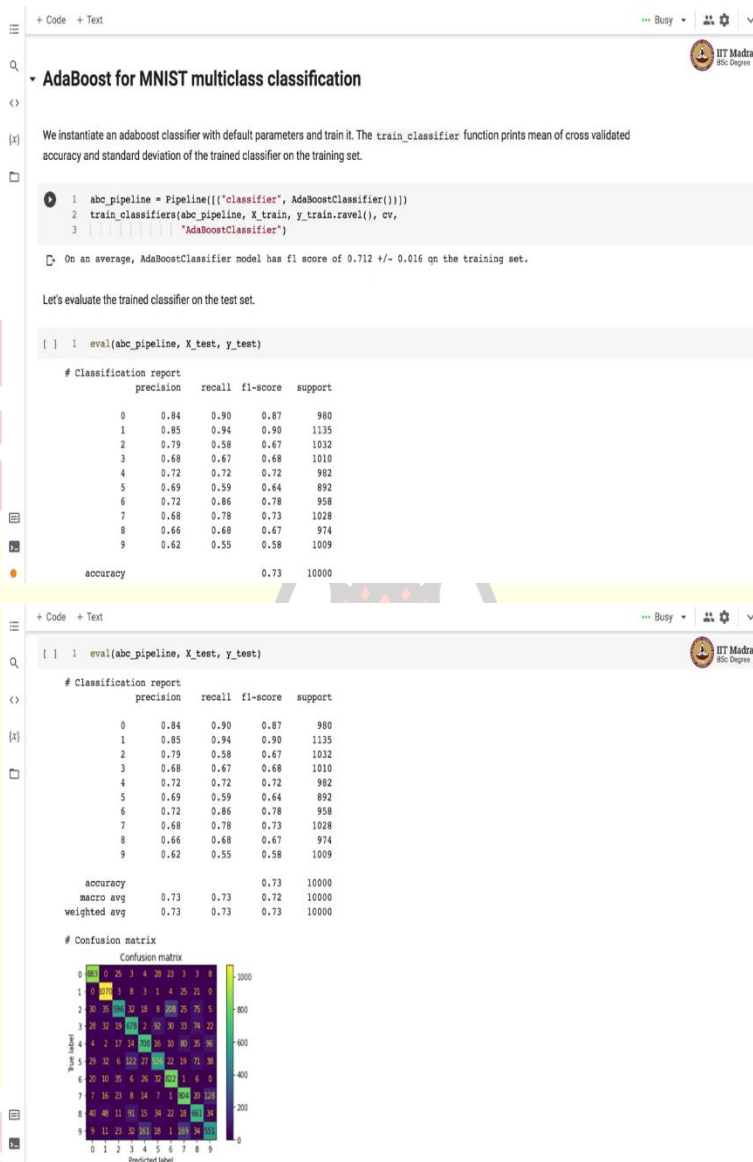
Let's evaluate the trained classifier on the test set.

```
1 eval(abc_pipeline, X_test, y_test)
```

# Classification report				
	precision	recall	f1-score	support
0	0.84	0.90	0.87	980
1	0.85	0.94	0.90	1135

We will train 2 classifiers with default parameters 1 is AdaBoostClassifier, and 2 is GradientBoostingClassifier. AdaBoostClassifier uses decision stumps as weak learner whereas GradientBoostingClassifier uses decision trees as weak learners.

(Refer Slide Time: 03:56)

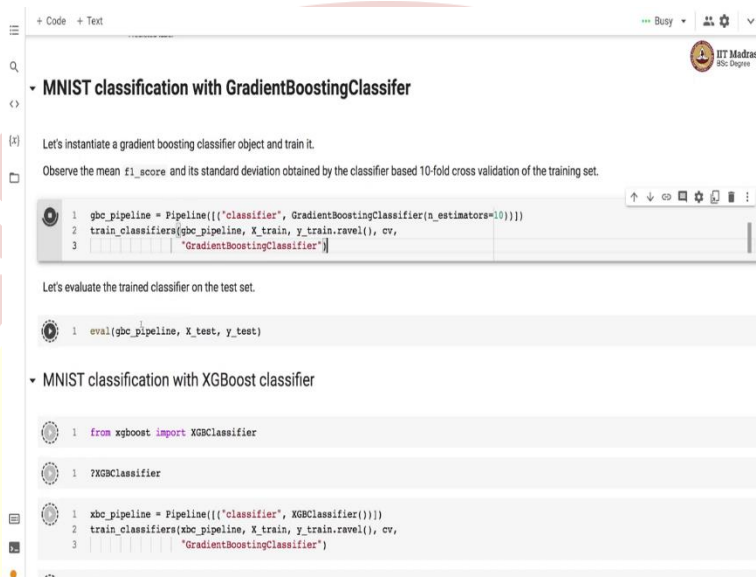


We instantiate `AdaBoostClassifier` with default parameters and train it with `train_classifier` function. The `train_classifier` function prints mean of cross validated accuracy and standard deviation of the train classifier on the training set. Here, we see that `AdaBoostClassifier` model obtains f1\_score of 0.712 with a standard deviation of 0.016 on the training set. And you can see that `AdaBoostClassifier` does not really get us that great classifier if you compare this with random forest classifier or bagging classifier that we saw in the previous collab.



So, there are a lot of confusions that are happening between different classes, for example, class 9 and 4 or class 7 and class 9, or class 5 and class 3 or class 6 and class 2, there are a lot of confusions and because of which the accuracy is merely 0.73 or 73%.

**(Refer Slide Time: 04:59)**



```
+ Code + Text
Busy
IIT Madras
PG Diploma

MNIST classification with GradientBoostingClassifier

Let's instantiate a gradient boosting classifier object and train it.
Observe the mean f1_score and its standard deviation obtained by the classifier based 10-fold cross validation of the training set.

1 gbc_pipeline = Pipeline([("classifier", GradientBoostingClassifier(n_estimators=10))])
2 train_classifiers(gbc_pipeline, X_train, y_train.ravel(), cv,
3                  "GradientBoostingClassifier")

Let's evaluate the trained classifier on the test set.

1 eval(gbc_pipeline, X_test, y_test)

MNIST classification with XGBoost classifier

1 from xgboost import XGBClassifier

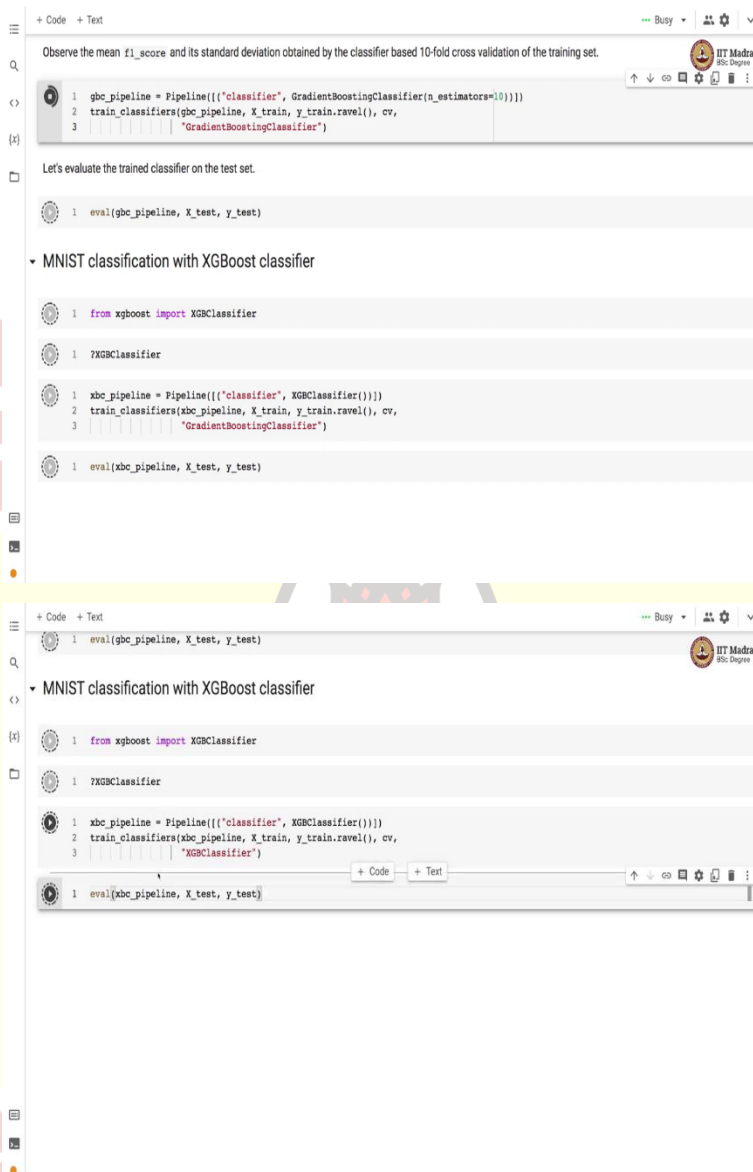
2 XGBClassifier

1 xbc_pipeline = Pipeline([("classifier", XGBClassifier())])
2 train_classifiers(xbc_pipeline, X_train, y_train.ravel(), cv,
3                  "GradientBoostingClassifier")
```

Now here we have given the code for training the model with GradientBoostingClassifier and XGBoost classifier, and as an exercise, you should run this code and obtain the accuracy with the GradientBoostingClassifier as well as XGBoost classifier and compare it with the accuracy that you obtained through AdaBoostClassifier.

So, here what we have done is we have created a pipeline object with GradientBoostingClassifier as 1 of the stages and the GradientBoostingClassifier is instantiated with number of estimators = 10. And then we train the GradientBoostingClassifier pipeline with train\_classifier function. So, what you have to do is you have to run this and find out what is the performance that we get and also run the evaluation pipeline and find out the performance on the test set from the GradientBoostingClassifier.

(Refer Slide Time: 06:01)



The image displays two screenshots of a Jupyter Notebook interface. The top screenshot shows the training and evaluation of a Gradient Boosting Classifier (gbc\_pipeline). The code includes importing GradientBoostingClassifier, creating a pipeline, training on training data, and evaluating on test data. The bottom screenshot shows the training and evaluation of an XGBClassifier (xgc\_pipeline). The code includes importing XGBClassifier, creating a pipeline, training on training data, and evaluating on test data. Both screenshots show the Jupyter Notebook interface with a code editor and a console output area.

```
Observe the mean f1_score and its standard deviation obtained by the classifier based 10-fold cross validation of the training set.
```

```
1 gbc_pipeline = Pipeline([("classifier", GradientBoostingClassifier(n_estimators=10))])
2 train_classifiers(gbc_pipeline, X_train, y_train.ravel(), cv,
3                   "GradientBoostingClassifier")
```

Let's evaluate the trained classifier on the test set.

```
1 eval(gbc_pipeline, X_test, y_test)
```

▼ MNIST classification with XGBoost classifier

```
1 from xgboost import XGBClassifier
2 XGBClassifier
3 xgc_pipeline = Pipeline([("classifier", XGBClassifier())])
4 train_classifiers(xgc_pipeline, X_train, y_train.ravel(), cv,
5                   "XGBClassifier")
6 eval(xgc_pipeline, X_test, y_test)
```

```
1 eval(gbc_pipeline, X_test, y_test)
```

▼ MNIST classification with XGBoost classifier

```
1 from xgboost import XGBClassifier
2 XGBClassifier
3 xgc_pipeline = Pipeline([("classifier", XGBClassifier())])
4 train_classifiers(xgc_pipeline, X_train, y_train.ravel(), cv,
5                   "XGBClassifier")
6 eval(xgc_pipeline, X_test, y_test)
```

So, GradientBoostingClassifier generally takes longer to run than AdaBoostClassifier. XGBoost classifier on the other hand, is expected to be efficient version of GradientBoostingClassifier. It is very similar to GradientBoostingClassifier except that it performs regularization to obtain better generalization performance.

We can import the XGBClassifier from XGBoost module. So, XGBClassifier implements classification with XGBoost. If you are interested in reading about documentation now, XGBClassifier and you can access the documentation by putting the ? followed by XGBClassifier.

And if you run this cell, you will get to see the documentation of this classifier. So, here we instantiate a pipeline object with `XGBClassifier` as a classification stage. And we train the `XGBClassifier` with `train_classifier` function by supplying the estimator object, the feature matrix label vector, the cross-validation strategy and name of the classifier.

So, just like gradient boosting, you also have to run this particular code and check out what is the accuracy that or what is the f1-score that we get on the training set, as well as find out the accuracy on the test set. And as an exercise compare the accuracies of AdaBoost gradient boosting and XGBoost classifier.

So, with this, we have demonstrated how to perform classification with boosting techniques. Now you have three more classifiers in your tool set, which is `AdaBoostClassifier`, `GradientBoostingClassifier` and `XGBoost` classifier, `GradientBoostingClassifier` and `XGBoost` classifiers work well on structured data. In many of the Kaggle competitions, we have seen that `XGBoost` and `GradientBoostingClassifier` give the state of the art performance whenever they are applied on structured dataset.

